



Visual Motion & Motion Field & Optical flow estimation

Motion Field

Motion Field : 2D projections of 3D displacement vectors due to camera and/or object motion




Consider a 3D point P and its image



Using pinhole camera equation:

$$p = \frac{fP}{Z}$$

Relative motion



The relative velocity of P wrt camera:

$$V = -t - \omega \times P$$


Translation
velocity

Rotation
angular
velocity

$$t = \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix}$$

$$\omega = \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

3D Relative Velocity



The relative velocity of P wrt camera:

$$V = -t - \omega \times P$$


$$t = \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} \quad \omega = \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix} \quad P = \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

$$V_x = -t_x - \omega_y Z + \omega_z Y$$

$$V_y = -t_y - \omega_z X + \omega_x Z$$

$$V_z = -t_z - \omega_x Y + \omega_y X$$

Motion Field



$$p = \frac{fP}{Z}$$


Taking derivative wrt time:

$$\begin{aligned} \frac{dp}{dt} = \mathbf{v} &= \frac{d \frac{fP}{Z}}{dt} = \frac{f}{Z^2} \left[\frac{dP}{dt} \cdot Z - P \cdot \frac{dZ}{dt} \right] \\ &= \frac{f}{Z^2} [V \cdot Z - P \cdot V_z] \end{aligned}$$

the velocity of p

$$\mathbf{v} = f \frac{V}{Z} - p \frac{V_z}{Z}$$

Motion Field




$$v = f \frac{V}{Z} - p \frac{V_z}{Z}$$

$$v_x = u = f \frac{V_x}{Z} - x \frac{V_z}{Z}$$

$$v_y = v = f \frac{V_y}{Z} - y \frac{V_z}{Z}$$

$$v_z = f \frac{V_z}{Z} - f \frac{V_z}{Z} = 0$$

Motion Field



$$u = f \frac{V_x}{Z} - x \frac{V_z}{Z}$$

$$v = f \frac{V_y}{Z} - y \frac{V_z}{Z}$$

$$V_x = -T_x - \omega_y Z + \omega_z Y$$


$$V_y = -T_y - \omega_z X + \omega_x Z$$

$$V_z = -T_z - \omega_x Y + \omega_y X$$

$$u = \frac{t_z x - t_x f}{Z} + \frac{\omega_x x y}{f} - \omega_y \left(f + \frac{x^2}{f} \right) + \omega_z y$$

$$v = \frac{t_z y - t_y f}{Z} + \omega_x \left(f + \frac{y^2}{f} \right) - \frac{\omega_y x y}{f} - \omega_z x$$

Motion Field:




Translational
component

Scaling ambiguity
(t and Z can only
be derived up to a scale
Factor)

$$u = \frac{t_z x - t_x f}{Z} + \frac{\omega_x x y}{f} - \omega_y \left(f + \frac{x^2}{f} \right) + \omega_z y$$

$$v = \frac{t_z y - t_y f}{Z} + \omega_x \left(f + \frac{y^2}{f} \right) - \frac{\omega_y x y}{f} - \omega_z x$$

Motion Field:




Rotational
component

$$\begin{aligned}
 u &= \frac{t_z x - t_x f}{Z} + \frac{\omega_x x y}{f} - \omega_y \left(f + \frac{x^2}{f} \right) + \omega_z y \\
 v &= \frac{t_z y - t_y f}{Z} + \omega_x \left(f + \frac{y^2}{f} \right) - \frac{\omega_y x y}{f} - \omega_z x
 \end{aligned}$$

NOTE: The rotational component is independent of depth Z !

Motion Field



$$\mathbf{v}(x, y) = \frac{1}{Z(x, y)} \mathbf{A}(x, y) \mathbf{V} + \mathbf{B}(x, y) \boldsymbol{\omega}$$

$$\mathbf{v}(x, y) = \begin{bmatrix} u(x, y) \\ v(x, y) \end{bmatrix}$$

$$\mathbf{A}(x, y) = \begin{bmatrix} -f & 0 & x \\ 0 & -f & y \end{bmatrix}$$

$$\mathbf{B}(x, y) = \begin{bmatrix} \frac{xy}{f} & -f - \frac{x^2}{f} & y \\ f + \frac{y^2}{f} & -\frac{xy}{f} & x \end{bmatrix}$$

$$\mathbf{V} = \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix}$$


$$\boldsymbol{\omega} = \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

Translational field

$$t_z \neq 0$$

$$u_{tr}(x,y) = (x - x_o) \frac{t_z}{Z}$$

$$v_{tr}(x,y) = (y - y_o) \frac{t_z}{Z}$$



where $p_o = (x_o, y_o) = \left(\frac{t_x}{t_z} \cdot f, \frac{t_y}{t_z} \cdot f \right)$ is the focus of expansion (FOE)


or focus of contraction (FOC).

Translational field

$$t_z = 0$$

$$u_{tr}(x, y) = -\frac{t_x \cdot f}{Z}$$

$$v_{tr}(x, y) = -\frac{t_y \cdot f}{Z}$$



where $p_o = (x_0, y_0) = \left(\infty, \infty \right)$ is the focus of expansion (FOE)

or focus of contraction (FOC).

All motion field vectors are parallel to each other and inversely proportional to depth !

Pure Translation: Properties of the MF

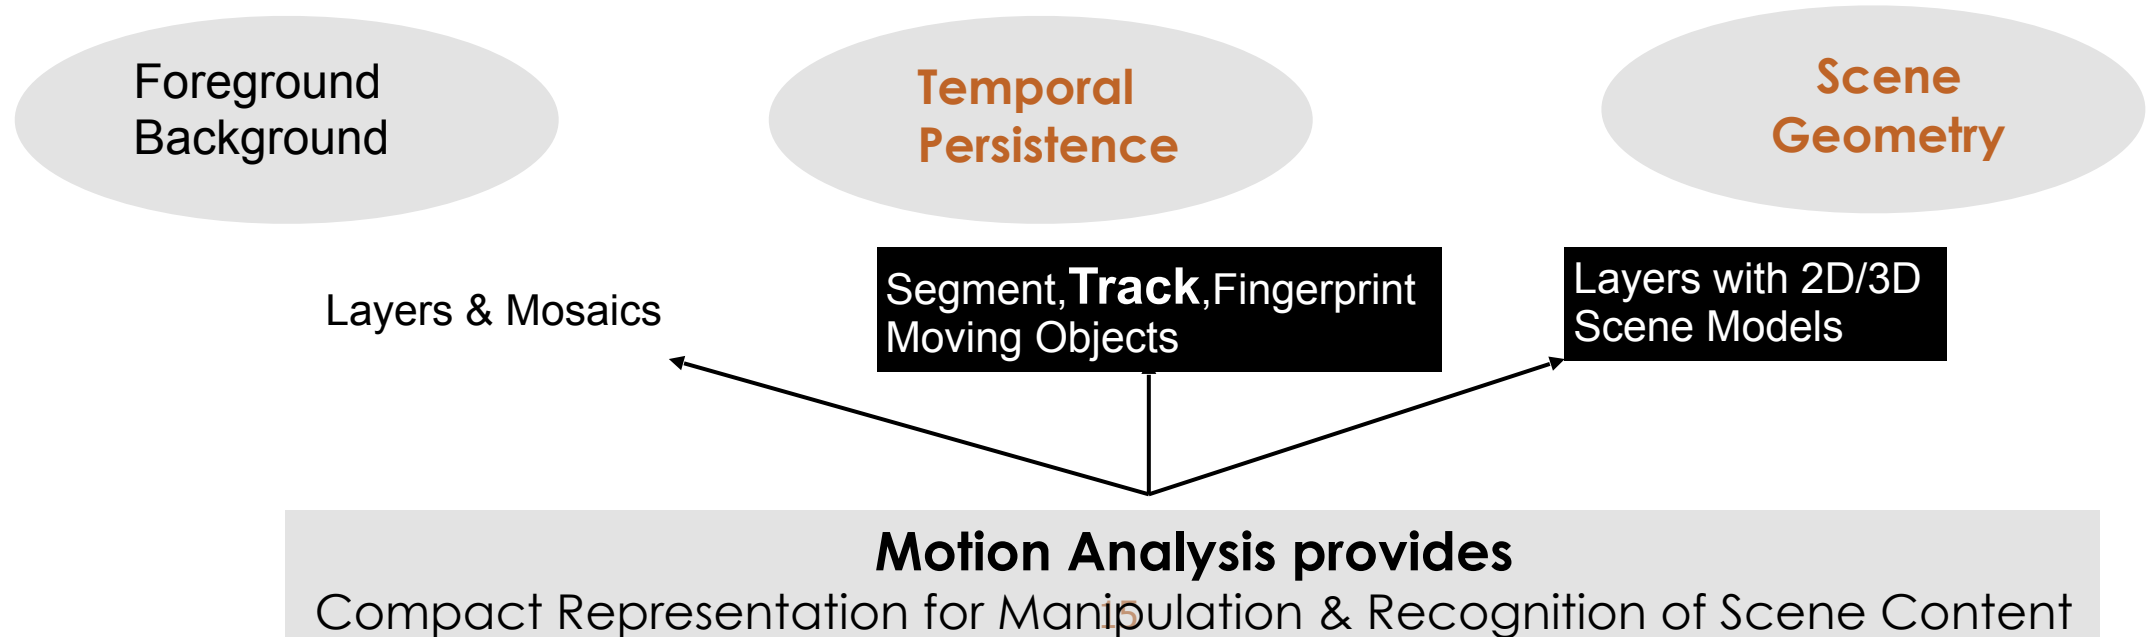
- *If $t_z \neq 0$ the MF is RADIAL with all vectors pointing towards (or away from) a single point p_o . If $t_z = 0$ the MF is PARALLEL.*
- *The length of the MF vectors is inversely proportional to depth Z . If $t_z \neq 0$ it is also directly proportional to the distance between p and p_o .*

$$u_{tr}(x,y) = (x - x_o) \frac{t_z}{Z}$$
$$v_{tr}(x,y) = (y - y_o) \frac{t_z}{Z}$$


Motion Analysis

$$\mathbf{v}(x, y) = \frac{1}{Z(x, y)} \mathbf{A}(x, y) \mathbf{V} + \mathbf{B}(x, y) \omega$$


$$\mathbf{V} \Rightarrow \{Z, \mathbf{V}, \omega\}$$




Typical Motion Fields




Zoom out




Zoom in




Pan right to left




Forward motion



Rotation



Horizontal translation



Closer objects appear to move faster!!

Optical Flow Estimation

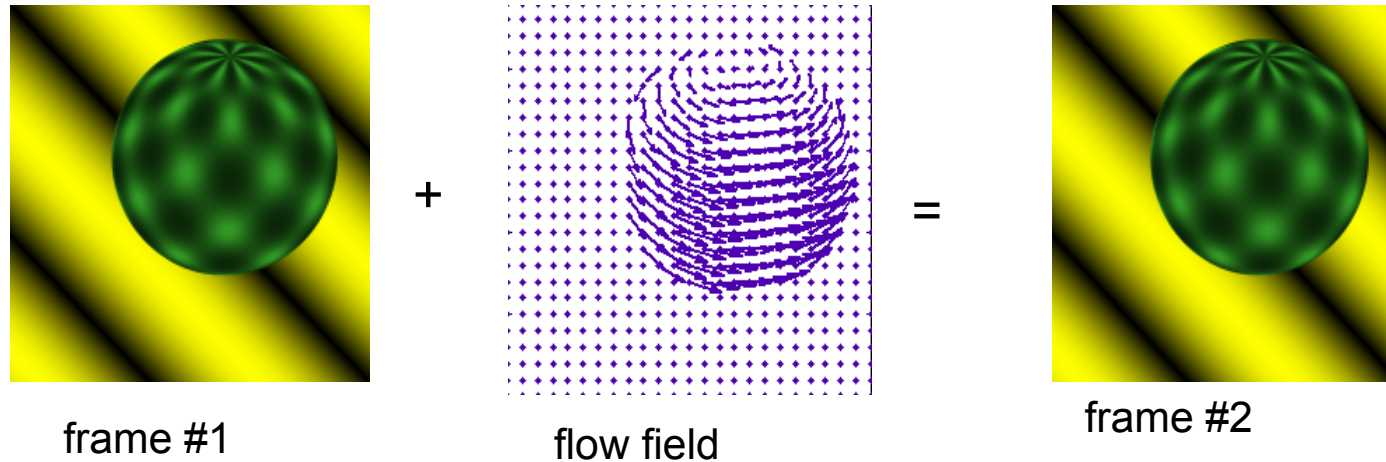
- Optical Flow
 - *Brightness constancy constraint*
 - *Aperture problem*
 - *Lucas-Kanade flow*
 - *Iterative refinement*
 - *Coarse-to-fine estimation*
- Global parametric motion
- Global Optical flow Constraint
- GFeature Tracking (sparse optical flow)

- **Definition-1: optical flow is the apparent motion of brightness patterns in the image.**
- **Ideally**, optical flow would be the same as the motion field.
- Have to be careful: **apparent motion can be caused by lighting changes without any actual motion.**
 - Think of a uniform rotating sphere under fixed lighting vs. a stationary sphere under moving illumination.

Optical Flow & Motion Field


■ Definition-2

The **optical flow** is a **velocity field** in the image which transforms one image into the next image in a sequence
[Horn&Schunck]




The **motion field** ... is the projection into the image of three-dimensional motion vectors [Horn&Schunck]


Optical Flow & Motion Field



Estimating Optical Flow




$I(x,y,t-1)$



$I(x,y,t)$

- Given two subsequent frames, estimate the apparent motion field $u(x,y)$ and $v(x,y)$ between them.
- **Key assumptions**
 - **Brightness constancy**: projection of the same point looks the same in every frame. (**Local image constraints**)
 - **Small motion**: points do not move very far.
 - **Spatial coherence**: points move like their neighbors.

Local image constraints



- **Brightness Constancy Equation:**

$$I(x, y, t-1) = I(x + u(x, y), y + v(x, y), t)$$

- **Linearizing the right hand side using Taylor expansion:**

$$I(x, y, t-1) \approx I(x, y, t) + I_x \cdot u(x, y) + I_y \cdot v(x, y)$$

- Hence, $I_x \cdot u + I_y \cdot v + I_t \approx 0$

Spatial derivatives

Temporal derivative

Local image constraints

$$I_x \cdot u + I_y \cdot v + I_t = 0$$


- How many equations and unknowns per pixel?
 - One equation, two unknowns

- Intuitively, what does this constraint mean?

$$\nabla I \cdot (u, v) + I_t = 0$$

- The component of the flow perpendicular to the gradient (i.e., parallel to the edge) is unknown

If (u, v) satisfies the equation,
so does $(u+u', v+v')$ if $\nabla I \cdot (u', v') = 0$



Solving the Aperture Problem

- How to get more equations for a pixel?
- **Spatial coherence constraint:** pretend the pixel's neighbors have the same (u,v)
 - If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(p_i) + \nabla I(p_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

B. Lucas and T. Kanade. **An iterative image registration technique with an application to stereo vision**. In Proceedings of the International Joint Conference on Artificial Intelligence, pp. 674–679, 1981

- **Least squares problem:**

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix} \quad \begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{matrix}$$

- **Minimum least squares solution given by solution of**

$$\begin{matrix} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 & 2 \times 1 \end{matrix}$$


$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$$\begin{matrix} A^T A & A^T b \end{matrix}$$

(The summations are over all pixels in the K x K window)

Interpreting the Eigenvalues

- Classification of image points using eigenvalues of the second moment matrix:



Iterative Refinement

1. Estimate velocity at each pixel using one iteration of Lucas and Kanade estimation.

$$\underbrace{\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix}}_{A^T A} \begin{bmatrix} u \\ v \end{bmatrix} = - \underbrace{\begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}}_{A^T b}$$

2. Warp one image toward the other using the estimated flow field.
3. Refine estimate by repeating the process.

B. Lucas and T. Kanade. **An iterative image registration technique with an application to stereo vision**. In Proceedings of the International Joint Conference on Artificial Intelligence, pp. 674–679, 1981

Lucas-Kanade flow

- Brightness constant equation (Optical equation)

$$I(x, y, t - 1) = I(x + u(x, y), y + v(x, y), t)$$

- Spatial coherence constraint (local): pretend the pixel's neighbors (Ω) have the same (u, v)

$$E(u, v) = \sum_{(x, y) \in \Omega} w(x, y) (I_x(x, y) \cdot u + I_y(x, y) \cdot v + I(x, y, t) - I(x, y, t - 1))^2$$

$$= \sum_{(x, y) \in \Omega} w(x, y) \left(\begin{bmatrix} I_x & I_y & I_t \end{bmatrix} \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \right)^2 = \begin{bmatrix} u & v & 1 \end{bmatrix} M \begin{bmatrix} u \\ v \\ 1 \end{bmatrix}$$

$$M = \sum_{(x, y) \in \Omega} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y & I_x I_t \\ I_y I_x & I_y^2 & I_y I_t \\ I_t I_x & I_t I_y & I_t^2 \end{bmatrix}$$

Lucas-Kanade flow

- Solve independently for each point [Lucas & Kanade 1981]

$$\frac{\partial E(u, v)}{\partial (u, v)} = 0 \Rightarrow \begin{bmatrix} \sum_{(x, y) \in \Omega} w(x, y) I_x^2 & \sum_{(x, y) \in \Omega} w(x, y) I_x I_y \\ \sum_{(x, y) \in \Omega} w(x, y) I_y I_x & \sum_{(x, y) \in \Omega} w(x, y) I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum_{(x, y) \in \Omega} w(x, y) I_x I_t \\ \sum_{(x, y) \in \Omega} w(x, y) I_y I_t \end{bmatrix}$$

$$G_\sigma * \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = -G_\sigma * \begin{bmatrix} I_x I_t \\ I_y I_t \end{bmatrix}$$

Affine Motion


- Affine Motion

$$u(x, y) = a_1 + a_2x + a_3y$$

$$v(x, y) = a_4 + a_5x + a_6y$$

- Substituting into the brightness constancy equation:

$$I_x \cdot u + I_y \cdot v + I_t \approx 0$$



Affine Motion

- Affine Motion

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

- Substituting into the brightness constancy equation:

$$I_x(a_1 + a_2 x + a_3 y) + I_y(a_4 + a_5 x + a_6 y) + I_t \approx 0$$

- Each pixel provides 1 linear constraint in 6 unknowns.
- Spatial coherence constrains, Least squares minimization:


$$Err(\vec{a}) = \sum \left[I_x(a_1 + a_2 x + a_3 y) + I_y(a_4 + a_5 x + a_6 y) + I_t \right]^2$$

Problem Cases in Lucas-Kanade

- **The motion is large** (larger than a pixel)
 - *Iterative refinement, coarse-to-fine estimation*
- **A point does not move like its neighbors**
 - *Motion segmentation*
- **Brightness constancy does not hold**
 - *Do exhaustive neighborhood search with **normalized correlation**.*

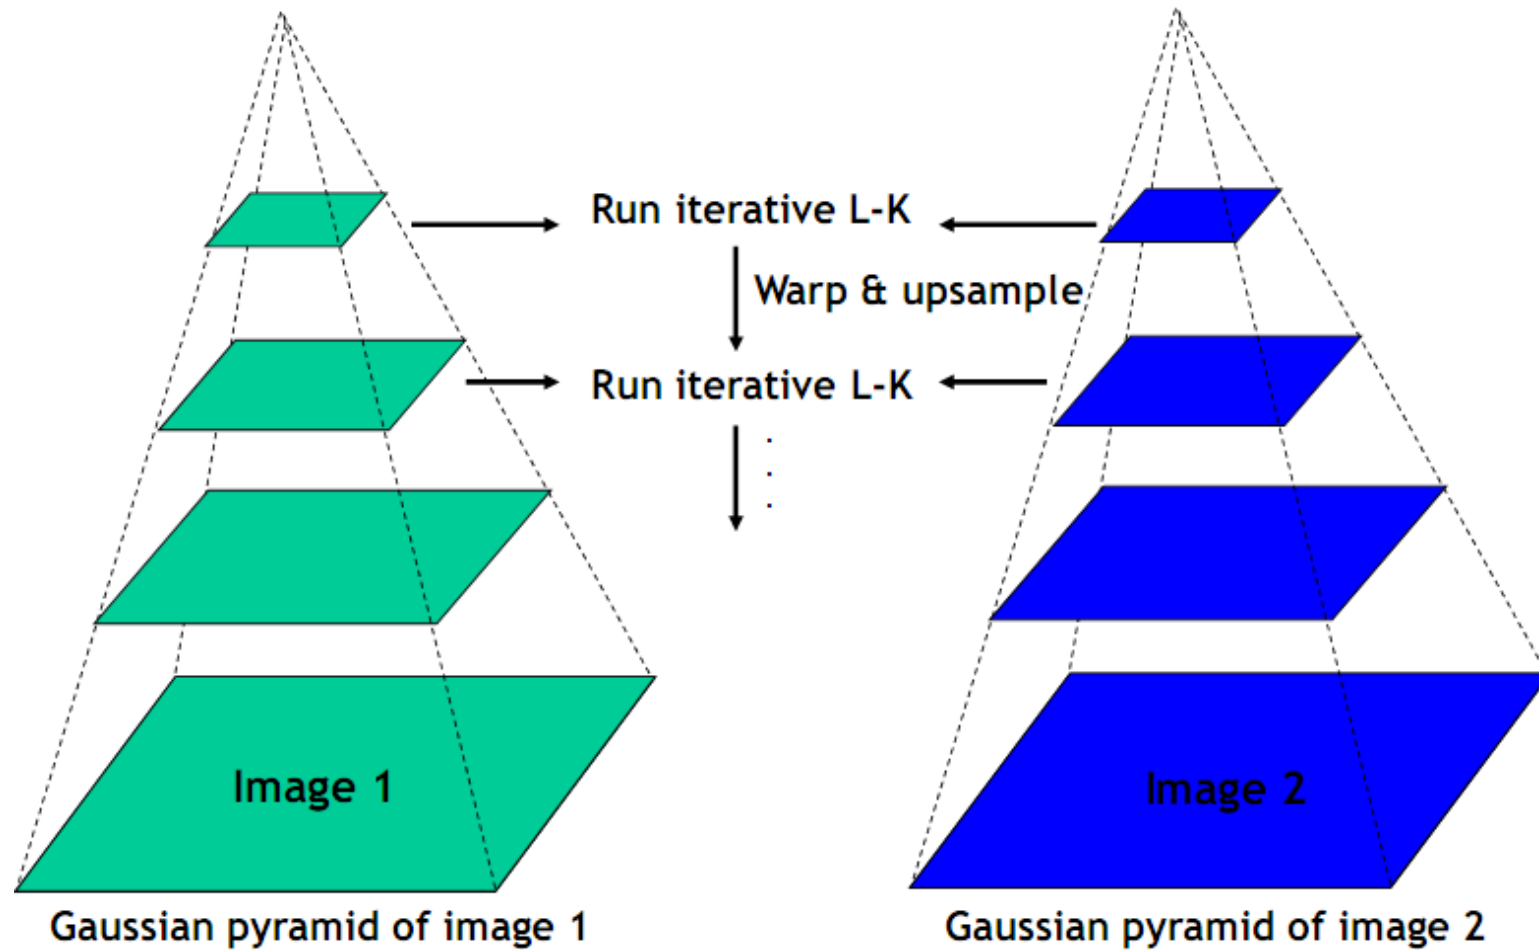
Dealing with Large Motions/ Temporal Aliasing

- Temporal aliasing causes ambiguities in optical flow because images can have **many pixels with the same intensity**.
 - I.e., how do we know which 'correspondence' is correct?



To overcome aliasing: coarse-to-fine estimation.

Coarse-to-fine Optical Flow Estimation



Jean-Yves Bouguet, *Pyramidal Implementation of the Lucas Kanade Feature Tracker*, TR, Intel, , 1997

Extension: Gradient constancy

Brightness is not always constant



Rotating cylinder


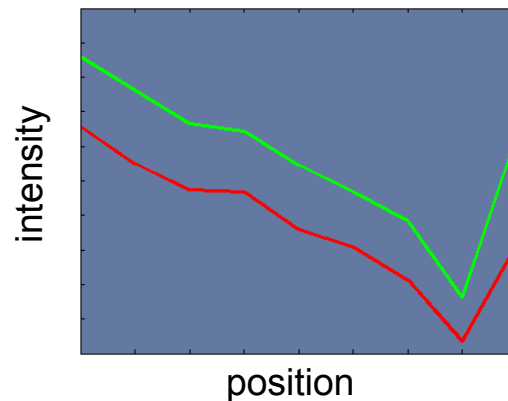


Brightness constancy does not always hold

Gradient constancy holds

$$I(x + u, y + v, t + 1) \neq I(x, y, t)$$

$$\nabla I(x + u, y + v, t + 1) = \nabla I(x, y, t)$$



Local constraints (data) + Local spatial coherence

➤ Brightness constancy

$$I(x + u, y + v, t + 1) - I(x, y, t) = 0$$

linearized

$$[u \ v \ 1] J \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \cong 0$$

Local spatial coherence

averaged linearized

$$\delta_{LIN+GAUSS}^2 = [u \ v \ 1] (G_\rho * J) \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \cong 0$$

➤ Gradient constancy

$$\nabla I(x + u, y + v, t + 1) - \nabla I(x, y, t) = 0$$

Feature Tracking

(sparse optical flow)

Tracking Challenges

- **Ambiguity of optical flow**
 - *Find good features to track*
- **Large motions**
 - *Discrete search instead of Lucas-Kanade*
- **Changes in shape, orientation, color**
 - *Allow some matching flexibility*
- **Occlusions, disocclusions**
 - *Need mechanism for deleting, adding new features*
- **Drift** – errors may accumulate over time
 - *Need to know when to terminate a track*

Handling Large Displacements

- Define a small area around a pixel as the **template**.
- Match the template against each pixel **within a search area** in next image — *just like stereo matching!*
- Use a match measure such as SSD or correlation.
- After finding the best discrete location, can use Lucas-Kanade to get **sub-pixel estimate**.

离散粗定位+基于梯度精细搜索

Tracking Over Many Frames

- Select features in first frame
- For each frame:
 - Update positions of tracked features
 - **Discrete search** or **Lucas-Kanade** (Image gradient)
 - Terminate inconsistent tracks
 - **Compute similarity with corresponding feature** in the previous frame or in the first frame where it's visible
- Start new tracks if needed
 - Typically every ~ 10 frames, new features are added to “refill the ranks”

Shi-Tomasi Feature Tracker

- Find good features using eigenvalues of second-moment matrix
 - Key idea: “good” features to track are the ones that can be tracked reliably.
- From frame to frame, track with Lucas-Kanade and a pure translation model.
 - More robust for small displacements, can be estimated from smaller neighborhoods.
- Check consistency of tracks by affine registration to the first observed instance of the feature.
 - Affine model is more accurate for larger displacements.
 - Comparing to the first frame helps to minimize drift.

KLT—Pyramidal tracking algorithm (Coarse-to-fine)

Goal: Let \mathbf{u} be a point on image I . Find its corresponding location \mathbf{v} on image J

Build pyramid representations of I and J : $\{I^L\}_{L=0,\dots,L_m}$ and $\{J^L\}_{L=0,\dots,L_m}$

Initialization of pyramidal guess: $\mathbf{g}^{L_m} = [g_x^{L_m} \ g_x^{L_m}]^T = [0 \ 0]^T$

for $L = L_m$ **down to** 0 **with step of** -1

Location of point \mathbf{u} on image I^L : $\mathbf{u}^L = [p_x \ p_y]^T = \mathbf{u}/2^L$

Derivative of I^L with respect to x : $I_x(x, y) = \frac{I^L(x+1, y) - I^L(x-1, y)}{2}$

Derivative of I^L with respect to y : $I_y(x, y) = \frac{I^L(x, y+1) - I^L(x, y-1)}{2}$

Spatial gradient matrix:
$$G = \sum_{x=p_x-\omega_x}^{p_x+\omega_x} \sum_{y=p_y-\omega_y}^{p_y+\omega_y} \begin{bmatrix} I_x^2(x, y) & I_x(x, y) I_y(x, y) \\ I_x(x, y) I_y(x, y) & I_y^2(x, y) \end{bmatrix}$$

Initialization of iterative L-K: $\bar{\mathbf{v}}^0 = [0 \ 0]^T$

for $k = 1$ **to** K **with step of** 1 (or until $\|\bar{\eta}^k\| < \text{accuracy threshold}$)

Image warp

Image difference:

$$\delta I_k(x, y) = I^L(x, y) - J^L(x + g_x^L + \nu_x^{k-1}, y + g_y^L + \nu_y^{k-1})$$

Image mismatch vector:

$$\bar{b}_k = \sum_{x=p_x-\omega_x}^{p_x+\omega_x} \sum_{y=p_y-\omega_y}^{p_y+\omega_y} \begin{bmatrix} \delta I_k(x, y) I_x(x, y) \\ \delta I_k(x, y) I_y(x, y) \end{bmatrix}$$

Optical flow (Lucas-Kanade): $\bar{\eta}^k = G^{-1} \bar{b}_k$

Guess for next iteration: $\bar{v}^k = \bar{v}^{k-1} + \bar{\eta}^k$

end of for-loop on k

Final optical flow at level L : $\mathbf{d}^L = \bar{v}^K$

Guess for next level $L - 1$: $\mathbf{g}^{L-1} = [g_x^{L-1} \ g_y^{L-1}]^T = 2(\mathbf{g}^L + \mathbf{d}^L)$

end of for-loop on L

Final optical flow vector: $\mathbf{d} = \mathbf{g}^0 + \mathbf{d}^0$

Location of point on J : $\mathbf{v} = \mathbf{u} + \mathbf{d}$

Solution: The corresponding point is at location \mathbf{v} on image J

Real-Time GPU Implementations

- This basic **feature tracking framework** (Lucas-Kanade + Shi-Tomasi) is commonly referred to as “**KLT tracking**”.
 - Used as preprocessing step for many applications (recall the Structure-from-Motion pipeline)
 - Lends itself to easy parallelization
- Very fast GPU implementations available
 - C. Zach, D. Gallup, J.-M. Frahm, **Fast Gain-Adaptive KLT tracking on the GPU**. In CVGPU'08 Workshop, Anchorage, USA, 2008
 - 216 fps with automatic gain adaptation
 - 260 fps without gain adaptation

http://www.cs.unc.edu/~ssinha/Research/GPU_KLT/
<http://cs.unc.edu/~cmzach/opensource.html>

Summary

- ◆ Motion field & Optical flow field
- ◆ Optical flow equation & aperture problem
- ◆ LK's Method and Horn' Method
- ◆ Feature Tracking & Sparse optical flow